

Multi-core Programming

Lecture 07. Components of OpenMP

Today's Topics

- **What is OpenMP?**
- **The OpenMP Execution Model**
- **Components of OpenMP**
- **Refer to OpenMP 2.5 specification**
 - <http://www.openmp.org/mp-documents/spec25.pdf>

(也可是
3.0 版本 Specification)
3.1

What is OpenMP?

- Defacto standard API for writing *实际上成了标准* multithreaded programs in C/C++ and Fortran
 - Consists of
 - Compiler directives
 - Runtime routines
 - Environment Variables
- 构成(组成) [编译指令
运行库
环境变量]

A first OpenMP example

For-loop with independent iterations

```
for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
```

相互通信无须关心

Compiler directives
data-sharing

① Parallel for (没有指明如何分割)

② shared

③ private

④ 语义复杂。

For-loop parallelized using an OpenMP pragma

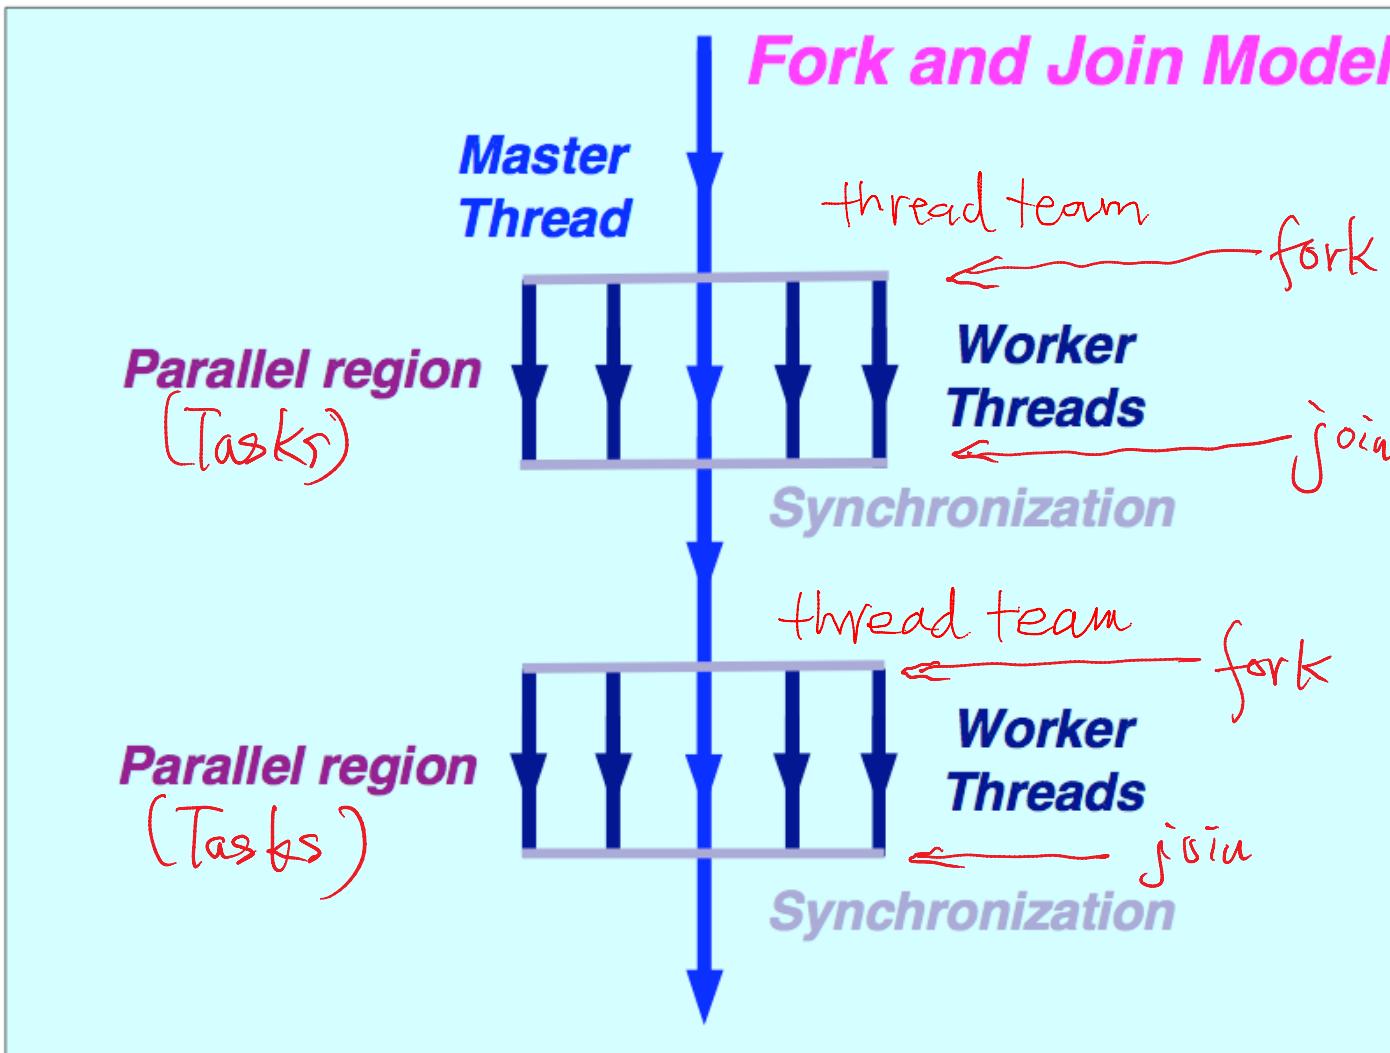
```
#pragma omp parallel for \
shared(n, a, b, c) \
private(i)
for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
```

并行化此 loop.
% cc -fopenmp source.c
% setenv OMP_NUM_THREADS 4
% a.out

环境变量

Fork-Join 扉絃技术

The OpenMP Execution Model



Terminology

- **OpenMP Team := Master + Workers** (Communicator in MPI)
- A Parallel Region is a block of code executed by all ~~线程区~~ threads simultaneously
 - The master thread always has thread ID 0 主线程
 - Thread adjustment (if enabled) is only done before entering a parallel region 并行区内线程调整操作
 - Parallel regions can be nested, but support for this is implementation dependent 并行区的嵌套
 - An "if" clause can be used to guard the parallel region; in case the condition evaluates to "false", the code is executed serially ~~线程区的guard~~ if语句的guard
- A work-sharing construct divides the execution of the enclosed code region among the members of the team; in other words: they split the work 工作共享结构实现工作分配

A first OpenMP example

For-loop with independent iterations

```
for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
```

Parallel Region
and Work-Sharing
Constructs

For-loop parallelized using an OpenMP pragma

```
#pragma omp parallel for \
shared(n, a, b, c) \
private(i)
for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
```

```
% cc -fopenmp source.c
% setenv OMP_NUM_THREADS 4
% a.out
```

Components of OpenMP

- Parallel regions
 - master & worker threads
 - Work-sharing constructs
 - parallel loop w/ schedule clauses
 - parallel sections
 - Single
 - Synchronization
 - nowait
 - Barrier
 - Data-sharing attributes
 - private, firstprivate, lastprivate, shared, reduction, threadprivate
 - Runtime routines
 - `omp_get_thread_num()`, `int omp_get_num_threads()`, `omp_in_parallel()`
 - Environment variables
 - `OMP_NUM_THREADS`, `OMP_SET_DYNAMIC`, `OMP_NESTED`, `OMP_SCHEDULE`
-
- The slide features several handwritten annotations in red ink:
- A large curly brace on the right side groups "Parallel regions", "Work-sharing constructs", and "Data-sharing attributes" under the heading "线程结构" (Thread Structure).
 - A curly brace on the right side groups "Synchronization" and "Runtime routines" under the heading "机制" (Mechanism).
 - Red text next to "Parallel regions" includes "并行区" (Parallel Region) and "工作共享结构" (Work-sharing Structure).
 - Red text next to "Synchronization" includes "同步机制" (Synchronization Mechanism).
 - Red text next to "Data-sharing attributes" includes "数据共享属性" (Data-sharing Attributes).
 - Red text next to "Runtime routines" includes "运行时函数" (Runtime Routines).
 - Red text next to "Environment variables" includes "环境变量" (Environment Variables).

OpenMP的构成以三种方式实现

Components of OpenMP

Directives 编译指令

- ◆ Parallel regions
- ◆ Work sharing
- ◆ Synchronization
- ◆ Data-sharing attributes
 - private
 - firstprivate
 - lastprivate
 - shared
 - reduction
- ◆ Orphaning

程序结构,同步,调度等

Environment variables 环境变量

- ◆ Number of threads
- ◆ Scheduling type
- ◆ Dynamic thread adjustment
- ◆ Nested parallelism

Runtime environment 运行时环境(线程)

- ◆ Number of threads
- ◆ Thread ID
- ◆ Dynamic thread adjustment
- ◆ Nested parallelism
- ◆ Timers
- ◆ API for locking

编译指令及语句

OpenMP directives and clauses

□ C: directives are case sensitive

- Syntax: #pragma omp directive [clause [clause] ...]

句法

□ Continuation: use \ in pragma

□ Conditional compilation: _OPENMP macro is set

if (scalar expression)

- Only execute in parallel if expression evaluates to true
- Otherwise, execute serially

private (list)

- No storage association with original object
- All references are to the local object
- Values are undefined on entry and exit

shared (list)

- Data is accessible by all threads in the team
- All threads access the same address space

```
#pragma omp parallel if (n > threshold) \
    shared(n,x,y) private(i)
{
    #pragma omp for
    for (i=0; i<n; i++)
        x[i] += y[i];
} /* End of parallel region */
```

if guard

编译指令子句
判断是否并行区

局部工作空间

firstprivate (list)

- All variables in the list are initialized with the value the original object had before entering the parallel construct

用原初值初始化

lastprivate (list)

- The thread that executes the sequentially last iteration or section updates the value of the objects in the list

从上一次串行段的结果更新值

Parallel Region

并行区

```
#pragma omp parallel [clause[,] clause] ...
{
    "this is executed in parallel" 代码
} (implied barrier) ← 暗含的全局同步
```

句法

A parallel region is a block of code executed by multiple threads simultaneously, and supports the following clauses:

| | | |
|--------------|---------------------|---------------------|
| if | (scalar expression) | 并行化条件 |
| private | (list) | 表明私有变量 |
| shared | (list) | 共享 |
| default | (nonshared) | 设置数据共享属性(C/C++) |
| default | (nonshared private) | (Fortran) |
| reduction | (operator: list) | 缩减操作 |
| copyin | (list) | master线程将数据从其他线程中读取 |
| firstprivate | (list) | 按原始值被初始化 |
| num_threads | (scalar_int_expr) | 线程数 |

子句

OpenMP Clause List

- The clause list is used to specify conditional parallelization, number of threads, and data handling.
 - Conditional Parallelization: The clause **if (scalar expression)** determines whether the parallel construct results in creation of threads. 条件并行
 - Degree of Concurrency: The clause **num_threads(integer expression)** specifies the number of threads that are created. 并行度
 - Data Handling: The clause **private (variable list)** indicates variables local to each thread. The clause **firstprivate (variable list)** is similar to the private, except values of variables are initialized to corresponding values before the parallel directive. The clause **shared (variable list)** indicates that variables are shared across all the threads.

数据处理(共享所有)

Work-sharing constructs in a Parallel Region

工作共享(分派)机制

```
#pragma omp for  
{  
    .... ①  
}
```

FOR

```
#pragma omp sections  
{  
    .... ②  
}
```

Section

```
#pragma omp single  
{  
    .... ③  
}
```

另有 master
(workshare) [3]

- The work is distributed over the threads
- Must be enclosed in a parallel region 包含在并行块之内
- Must be encountered by all threads in the team, or none at all
- No implied barrier on entry; implied barrier on exit (unless nowait is specified) (master not)
- A work-sharing construct does not launch any new threads (无创建)
- Shorthand syntax supported for parallel region with single work-sharing construct e.g.,

```
#pragma omp parallel  
#pragma omp for  
for (...)
```



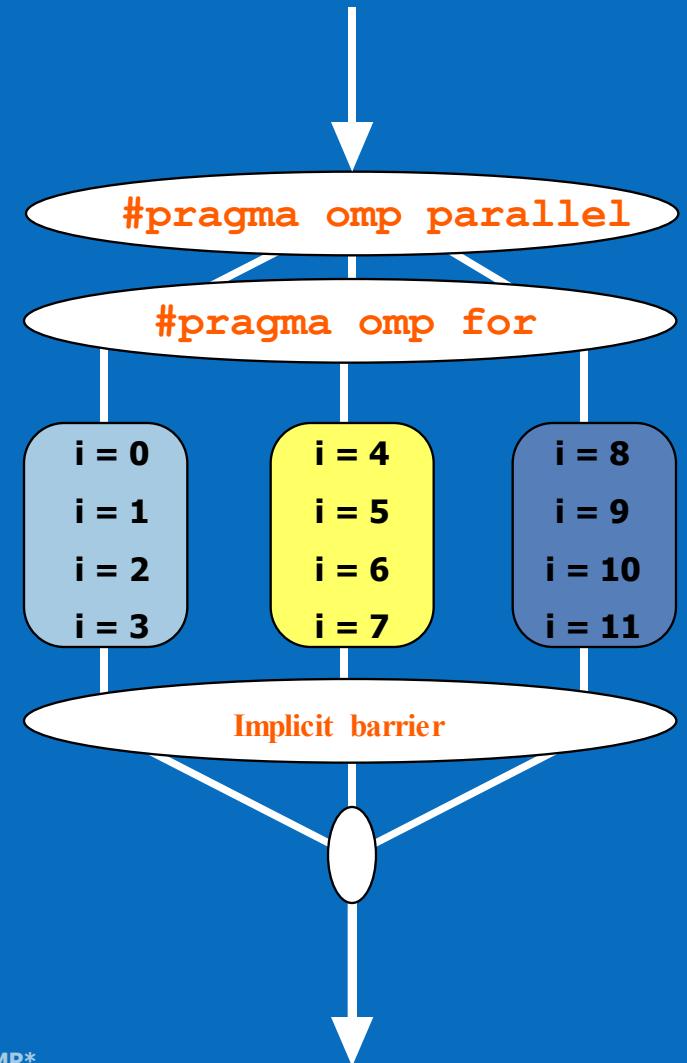
```
#pragma omp parallel for  
for (...)
```

Work-sharing Construct

```
#pragma omp parallel
#pragma omp for
for(i = 0; i < 12; i++)
    c[i] = a[i] + b[i]
```

Threads are assigned an independent set of iterations

Threads must wait at the end of work-sharing construct



Example of work-sharing “omp for” loop

```
#pragma omp parallel default(none) \
    shared(n,a,b,c,d) private(i)
{
    #pragma omp for nowait
    for (i=0; i<n-1; i++)  任务{派
        b[i] = (a[i] + a[i+1])/2;

    #pragma omp for nowait
    for (i=0; i<n; i++)  派
        d[i] = 1.0/c[i];
}

/*-- End of parallel region --*/
(implied barrier)
```



Reduction Clause in OpenMP

- The reduction clause specifies how multiple local copies of a variable at different threads are combined into a single copy at the master when threads exit.
- The usage of the reduction clause is **reduction (operator: variable list)**.
- The variables in the list are implicitly specified as being **private to threads**.
- The operator can be one of **+, *, -, &, |, ^, &&, and ||**.

```
#pragma omp parallel reduction(+ : sum) num_threads(8) {  
/* compute local sums here */  
}  
/*sum here contains sum of all local instances of sums */
```

OpenMP Programming: Example

计算π的并行程序

```
/* *****
```

An OpenMP version of a threaded program to compute PI.

```
***** */
```

```
#pragma omp parallel default(private) shared (npoints) \
reduction(+: sum) num_threads(8)
```

```
{
    num_threads = omp_get_num_threads();
    sample_points_per_thread = npoints / num_threads;
    sum = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        rand_no_x =(double)(rand_r(&seed))/(double)((2<<14)-1);
        rand_no_y =(double)(rand_r(&seed))/(double)((2<<14)-1);
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)
            sum++;
    }
}
```

of clause in h2?

Example of work-sharing “sections”

```
#pragma omp parallel default(none) \
    shared(n,a,b,c,d) private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        {
            for (i=0; i<n-1; i++)
                b[i] = (a[i] + a[i+1])/2;
        }
        #pragma omp section
        {
            for (i=0; i<n; i++)
                d[i] = 1.0/c[i];
        }
    } /*-- End of sections --*/
} /*-- End of parallel region --*/
```

用Section进行任务分配



“single” and “master” constructs in a parallel region

Only one thread in the team executes the code enclosed

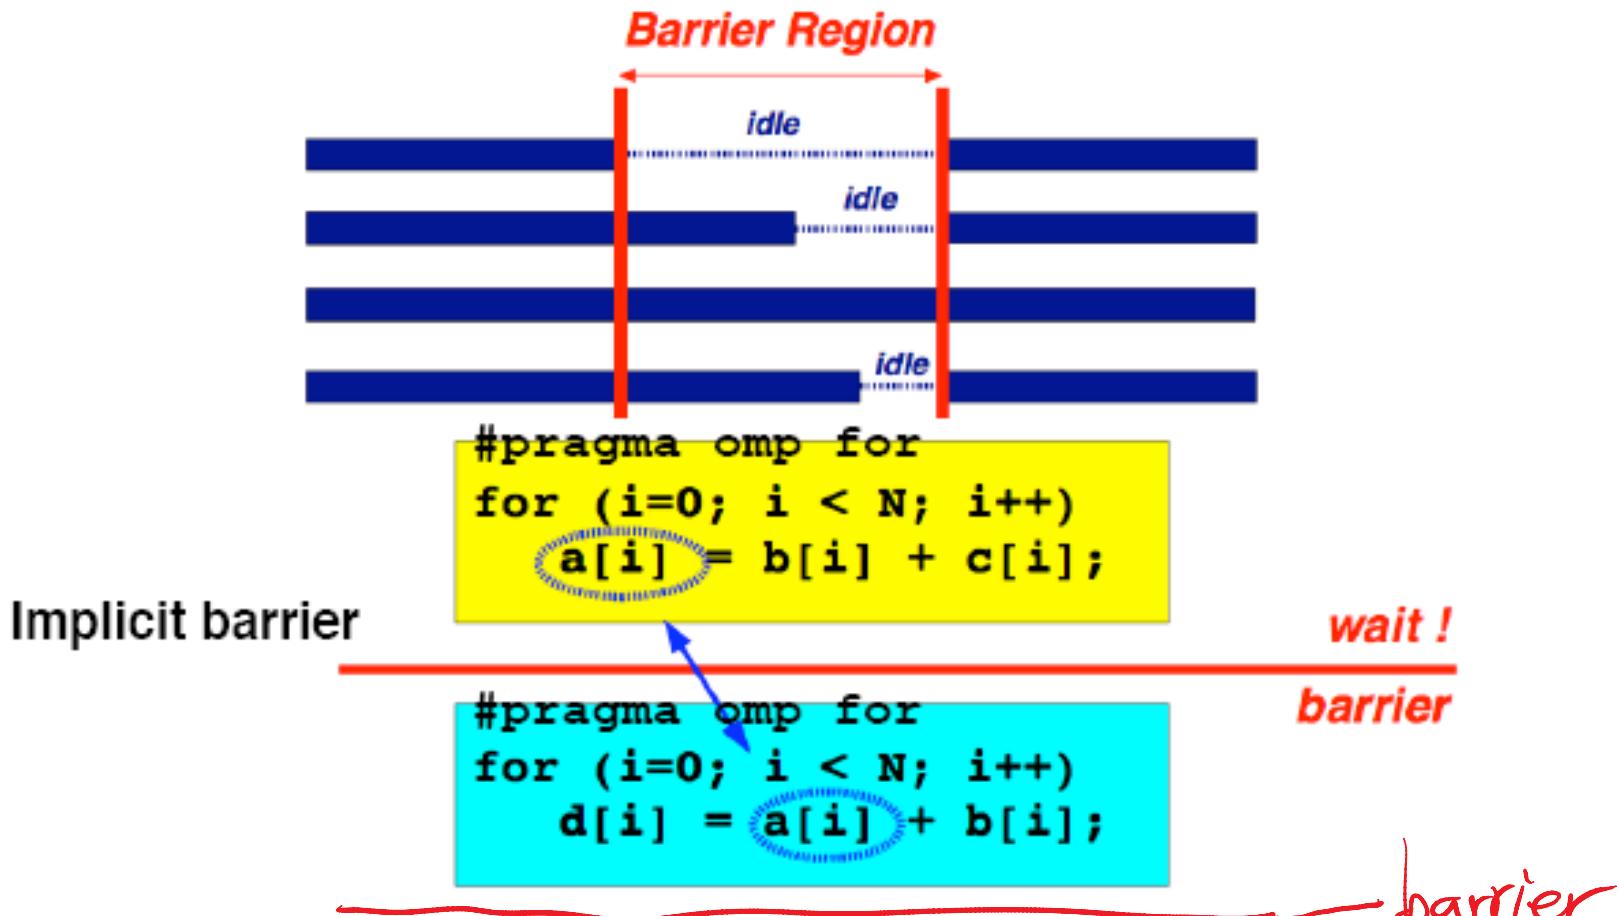
```
#pragma omp single [clause[,] clause] ...
{
    <code-block>
}
```

Only the master thread executes the code block,

```
#pragma omp master
{<code-block>}
```

- Single and master are useful for computations that are intended for single-processor execution e.g., I/O and initializations
- There is no implied barrier on entry or exit of a single or master construct
([↑] barrier)

Implicit barrier



NOTE: barrier is redundant if there is a guarantee that the mapping of iterations onto threads is identical in both loops

nowait clause & explicit barrier

```
#pragma omp for nowait  
{  
    :  
}
```

不等待 barrier

```
#pragma omp barrier
```

显式设置 barrier

- To minimize synchronization, some OpenMP directives/pragmas support the optional *nowait* clause
- If present, threads do not synchronize/wait at the end of that particular construct
- An explicit barrier can then be inserted at only the desired program points 在需要的地方插入显式 barrier.

A more elaborate example

```
#pragma omp parallel if (n>limit) default(none) \
    shared(n,a,b,c,x,y,z) private(f,i,scale)
{
    f = 1.0;
    #pragma omp for nowait
    for (i=0; i<n; i++)
        z[i] = x[i] + y[i];
    #pragma omp for nowait
    for (i=0; i<n; i++)
        a[i] = b[i] + c[i];
    #pragma omp barrier
    ...
    scale = sum(a,0,n) + sum(z,0,n) + f;
    ...
} /*-- End of parallel region --*/
```

parallel region

parallel loop
(work is distributed)

parallel loop
(work is distributed)

synchronization

Statement is executed by all threads

Statement is executed by all threads

The diagram illustrates the execution flow of the provided OpenMP code within a parallel region. It uses colored boxes and arrows to map specific sections of the code to annotations:

- Initial Setup:** The first section of code is annotated with "Statement is executed by all threads".
- Parallel Loop 1:** The first "#pragma omp for nowait" section is annotated with "parallel loop (work is distributed)".
- Parallel Loop 2:** The second "#pragma omp for nowait" section is annotated with "parallel loop (work is distributed)".
- Synchronization:** The "#pragma omp barrier" statement is annotated with "synchronization".
- Final Reduction:** The final section of code involving the calculation of "scale" is annotated with "Statement is executed by all threads".

Assigning Iterations to Threads

向线程分配迭代 (线程调度)

- The schedule clause of the **for** directive deals with the *assignment* of iterations to threads.
- The general form of the schedule directive is **schedule(scheduling_class[, parameter])**.
- OpenMP supports four scheduling classes: **static**, **dynamic**, **guided**, and **runtime**.

schedule clause for parallel loops

schedule (static | dynamic | guided [, chunk])
schedule (runtime)



static [, chunk]

块大小

- ✓ *Distribute iterations in blocks of size "chunk" over the threads in a round-robin fashion* 块大小分配{固定}
- ✓ *In absence of "chunk", each thread executes approx. N/P chunks for a loop of length N and P threads*

Example: Loop of length 16, 4 threads:

| TID | 0 | 1 | 2 | 3 |
|------------------|-------------|--------------|--------------|--------------|
| no chunk | 1-4 | 5-8 | 9-12 | 13-16 |
| chunk = 2 | 1-2 9-10 | 3-4 11-12 | 5-6 13-14 | 7-8 15-16 |

schedule clause for parallel loops (contd)

dynamic [, chunk]

- ✓ *Fixed portions of work; size is controlled by the value of chunk* 工作块大小固定, 线程动态调度, 每当一线程结束便开始下一块工作。
- ✓ *When a thread finishes, it starts on the next portion of work*

guided [, chunk]

同 dynamic, 但块大小按指数组递减

- ✓ *Same dynamic behavior as "dynamic", but size of the portion of work decreases exponentially* 指数组递减

runtime

- ✓ *Iteration scheduling scheme is set at runtime through environment variable OMP_SCHEDULE* 由环境变量决定

Assigning Iterations to Threads: Example

```
/* static scheduling of matrix multiplication loops */  
#pragma omp parallel default(private) shared (a, b, c, dim) \  
num_threads(4)  
#pragma omp for schedule(static)  
for (i = 0; i < dim; i++) {  
    for (j = 0; j < dim; j++) {  
        c(i,j) = 0;  
        for (k = 0; k < dim; k++) {  
            c(i,j) += a(i, k) * b(k, j);  
        }  
    }  
}
```

怎样分配？

Nesting parallel Directives

- Nested parallelism can be enabled using the **OMP_NESTED** environment variable.
- If the **OMP_NESTED** environment variable is set to **TRUE**, nested parallelism is enabled.
- In this case, each parallel directive creates a new team of threads. *每个指令将创建新线程.*

不在并行块之内 in parallel block 将被忽略.

Out-of-line (“orphaned”) directives

- ♦ *The OpenMP standard does not restrict worksharing and synchronization directives (`omp for`, `omp single`, `critical`, `barrier`, etc.) to be within the lexical extent of a parallel region. These directives can be orphaned*
- ♦ *That is, they can appear outside the lexical extent of a parallel region*

```
(void) dowork(); /* Sequential FOR  
#pragma omp parallel  
{  
    (void) dowork(); /* Parallel FOR  
}
```

```
void dowork()  
{  
#pragma omp for  
    for (i=0;....)  
    {  
        :  
    }  
}
```

- ♦ *When an orphaned worksharing or synchronization directive is encountered in the sequential part of the program (outside the dynamic extent of any parallel region), it is executed by the master thread only. In effect, the directive will be ignored*

Homework 6

- Read the section of *OpenMP C and C++ Application program Interface* in MSDN, and try to run some of the **examples** in this section.

或有任何地方找不到 OpenMP 程序, 请在 Linux
或 Windows 下 浏览运行。

(运行环境请自己安装配置)